

Lecture 16 MIP Solvers

Can Li

ChE 597: Computational Optimization
Purdue University

History of MIP solvers

- The CPLEX Optimizer was named for the simplex method implemented in the C programming language. It was originally developed by Robert Bixby and sold commercially from 1988 by CPLEX Optimization Inc. This was acquired by ILOG in 1997 and ILOG was subsequently acquired by IBM in January 2009.
- Zonghao Gu, Edward Rothberg, and Robert Bixby left CPLEX and founded Gurobi in 2008. Gurobi is similar to CPLEX by design. It was the best-performing MILP solver in the Mittelman benchmark. Now it has stopped participating in the Mittelman benchmark.
- Gurobi can now solve nonconvex MIQCQP. Starting from Gurobi 11, simple MINLP capability is added. Starting with Gurobi 12, general MINLP is added.


What's inside the box?

- Preprocessing
- Cuts
- Heuristics
- Branch and bound
- Callbacks

There are hundreds of parameters in CPLEX/Gurobi that control these procedures.

Preprocessing/presolve

Preprocessing routines can be run before the solving starts (usually several, sequentially), to simplify and / or tighten the problem formulation.



```
Presolve removed 83494 rows and 41814 columns  
Presolve time: 0.35s
```

One can control how much time/rounds to devote to preprocessing by changing the parameters. <https://www.gurobi.com/documentation/current/refman/presolve.html> to change a parameter in pyomo.

Cutting planes

Tighten the LP relaxation

- Boolean Quadric Polytope (BQP) cuts
- Clique cuts
- Cover cuts
- Disjunctive cuts
- Flow cover cuts
- Flow path cuts
- Gomory fractional cuts
- Zero-half cuts
- Generalized upper bound (GUB) cover cuts
- Implied bound cuts: global and local
- Lift-and-project cuts
- Mixed integer rounding (MIR) cuts
- Multi-commodity flow (MCF) cuts
- Reformulation Linearization Technique (RLT) cuts

There are parameters in CPLEX/Gurobi to control how “aggressive” you want to generate the cuts.

<https://docs.gurobi.com/projects/optimizer/en/current/reference/parameters.html#cuts>

Heuristics

Heuristics are to find good feasible solutions (upper bounds) quickly.

- Neighborhood search
- Local branching
- Relaxation induced neighborhood search (RINS)
- Solution polishing
- Feasibility pump

Parameter to control how much time to spend on heuristics

<https://www.gurobi.com/documentation/current/refman/heuristics.html>

Neighborhood search

Let x^* be the solution to the LP relaxation. Explore neighborhoods defined by $\lfloor x_i^* \rfloor \leq x_i \leq \lceil x_i^* \rceil, i \in I$, i.e., add these bound constraints to the MILP and solve the subMIP in the neighborhood of the LP solution. This only applies when x is has general integer variables. For binary variables, one can just do a simple rounding heuristic.

- The search within the subMIP is truncated by limiting the number of nodes explored.

Local Branching

- Assume we have a feasible solution \hat{x} , the so-called reference solution, and let $S := \{j \in I \mid \hat{x}_j = 1\}$ denote the binary support of \hat{x} .
 are the binary variables being 1.
 $I \setminus S$
- For a given positive integer parameter k , we define the k -OPT neighborhood $\mathcal{N}(\hat{x}, k)$ of \hat{x} as the set of the feasible solutions satisfying

$$\Delta(x, \hat{x}) := \sum_{j \in S} (1 - x_j) + \sum_{j \in I \setminus S} x_j \leq k, \quad (1)$$

$x \approx \hat{x}$

known as the *local branching constraint*.

- This constraint requires at most k variables have values different from \hat{x} .
- Constraint (1) can also be used to branch within a branch and bound:

$$\begin{aligned} \Delta(x, \hat{x}) &\leq k && \text{(left branch)} \\ \text{or } \Delta(x, \hat{x}) &\geq k + 1 && \text{(right branch)} \end{aligned}$$

Relaxation Induced Neighborhood Search

- A similar concept of neighborhood takes into account simultaneously both
 - the *incumbent solution* \bar{x} , and
 - the *solution of the continuous relaxation* \tilde{x} ,at a given node of the branch-and-bound tree.
- \bar{x} and \tilde{x} are compared and all the binary variables that assume the same value are **hard-fixed** in an associated MILP.
- This associated MILP is then solved by using the MILP solver as a black-box.
- In case the incumbent solution is improved, \bar{x} is updated in the rest of the tree.
- This method turns out to give very competitive results on general MILPs.

RINS Formulation

RINS: Let \tilde{x} be a known feasible solution and let \hat{x} be an LP-solution at some node in the search tree. We create a new MILP (Danna et al., 2005):

minimize $z = cx$
s.t.

$$\begin{aligned} Ax &\leq b, \\ x_i &= \tilde{x}_i \quad \forall i \quad \text{s.t. } \tilde{x}_i = \hat{x}_i \\ x &\in \mathbb{Z}^r \times \mathbb{R}^{n-r}. \end{aligned}$$

Solution Polishing


- Idea: Solution polishing can yield better solutions in situations where good solutions are otherwise hard to find. More time-intensive than other heuristics, solution polishing is actually a variety of branch and cut that works after an initial solution is available. In fact, it requires a solution to be available for polishing, either a solution produced by branch and cut, or a MIP start supplied by a user.
- Rothberg, 2007 developed an evolutionary algorithm for polishing MIP solutions. It can be seen as a combination of genetic algorithm (mutating, combining solutions from different MIPs). It requires solving sub MIPs.

Feasibility Pump: The Basic Scheme

- We start from any $x^0 \in P$ where P denotes the LP relaxation, and round to obtain \hat{x}^0 .
- We look for a point $\hat{x}^1 \in P$ which is as close as possible to \hat{x}^0 by solving the problem:

$$\min \{ \Delta(x, \hat{x}) \mid x \in P \}$$

distance metric.



- If we choose the measure $\Delta(x, \hat{x})$ properly, this problem is easily solvable.
- If \hat{x}^1 satisfies the integrality constraints, we are done.
- Otherwise, we obtain \hat{x}^1 by rounding \tilde{x}^1 , and repeat.
- From a geometric point of view, this is like an “alternating projection” algorithm (project to the polyhedron and project to the integer set). *converge when the two sets are convex*
- These satisfy feasibility in a complementary but partial way:
 1. \tilde{x}^i , satisfies the linear constraints,
 2. \hat{x}^i , the integrality requirements.

Branching Priority

- Affects the order of branching.
- For example, in scheduling problem, first branch on the assignment of jobs to machines, then branch on the sequence of the jobs.

<https://www.gurobi.com/documentation/current/refman/branchpriority.html>

Solution Pool

Goal: find multiple solutions in an MILP.

[https://www.gurobi.com/documentation/current/refman/
poolsearchmode.html](https://www.gurobi.com/documentation/current/refman/poolsearchmode.html)

Multi-objective

If you have multiple objectives $c_1^T x, c_2^T x, \dots, c_n^T x$, you can solve the problem

$$\min_x \sum_{i=1}^n \lambda_i c_i^T x$$

where $\lambda_i \geq 0, \forall i, \sum_{i=1}^n \lambda_i = 1$.

By changing the values of λ , we obtain the “Pareto front” of the multi-objective optimization problem where we cannot improve any objective without sacrificing others. https://www.gurobi.com/documentation/current/refman/multiple_objectives.html

LP algorithms

- Primal simplex
- Dual simplex
- Barrier with crossover
- Barrier without crossover

LP method <https://www.gurobi.com/documentation/current/refman/method.html>

Crossover <https://www.gurobi.com/documentation/current/refman/crossover.html>

- Takeaway: if you are solving an LP and doesn't care whether you get an extreme point or not, use "barrier without crossover" is fastest.
- If you are solving an MILP, the root relaxation has to return an extreme point. The reason is to use dual simplex to solve the node relaxations after branching starts. The dual feasible solution remain feasible after branching, which provides a good warmstart.

Parallel Computing

- By default, MIP solvers use all the threads on your computer/server. Branch and bound can be easily parallelized
- However, empirical results show it is best to use 6-8 threads. Having more threads may be counter-productive

<https://www.gurobi.com/documentation/current/refman/threads.html> Change the number of threads

Solver APIs

- The MIP solvers are usually written in C. But it has APIs in C, C++, Java, Python, Matlab, R. All these APIs are solver-dependent.

How does Pyomo work?

- Pyomo is designed to be a solver-independent algebraic modeling language.
- Pyomo can “talk” to different solvers. However, there are different ways to send your model data to the solvers.
 - LP file. write you model as a (.lp) file on your hard-drive. Then, the solver reads the LP file. This approach does not need to use the solver's python API.
 - Pyomo persistent solvers: communicate with the solver through the their python API.

Gurobi persistent solver:

```
https://github.com/Pyomo/pyomo/blob/main/pyomo/solvers/plugins/solvers/gurobi\_persistent.py
```

Lazy Constraints/Callback

- Lazy constraints are useful when the full set of constraints is too large to explicitly include in the initial formulation. When a MIP solver reaches a new solution, for example with a heuristic or by solving a problem at a node in the branch-and-bound tree, it will give the user the chance to provide constraints that would make the current solution infeasible.
- An example is the subtour elimination constraints.
- Lazy callbacks give more flexibility of what to implement when an integer feasible solution is found. This used in the implementation of single-tree Benders decomposition, single tree OA (QG).
- example
<https://stackoverflow.com/questions/58200552/pyomo-and-gurobi-does-pyomo-support-solver-callbacks-t>

User cut callback

- User cuts provide a way for the user to tighten the LP relaxation using problem-specific knowledge that the solver cannot or is unable to infer from the model. Just like with lazy constraints, when a MIP solver reaches a new node in the branch-and-bound tree, it will give the user the chance to provide cuts to make the current relaxed (fractional) solution infeasible in the hopes of obtaining an integer solution.

More on callbacks https://www.gurobi.com/documentation/current/refman/cb_codes.html

- Key take-away: Lazy constraints are necessary to get the optimal (feasible) solution. User cuts are not necessary (it only helps to solve the problem faster).

Solver log file

- **Read this carefully every time you solve a nontrivial optimization model!**
- Check matrix, rhs range. root node relaxation. root node bound after adding cuts.
- Easy problem can be solved fast, without doing much branch and bound.
- Hard problem takes exponentially many branch and bound nodes. Large memory consumption. Check your memory consumption from time to time or set an upper limit.
- Diagnose why it is slow: difficult to find good feasible solutions or closing the gap.

```
CPU model: AMD EPYC 7643 48-Core Processor, instruction set [SSE2|AVX|AVX2]
Thread count: 48 physical cores, 96 logical processors, using up to 8 threads

Optimize a model with 92153 rows, 50135 columns and 262513 nonzeros
Model fingerprint: 0x01f9d80d
Variable types: 29975 continuous, 20160 integer (20160 binary)
Coefficient statistics:
  Matrix range     [1e+02, 5e+02]
  Objective range  [9e+02, 2e+03]
  Bounds range     [0e+00, 0e+00]
  RHS range        [1e+00, 1e+08]
Found heuristic solution: objective -814886.9986
Presolve removed 83494 rows and 41814 columns
Presolve time: 0.35s
Presolved: 8661 rows, 8321 columns, 41716 nonzeros
Variable types: 6327 continuous, 1994 integer (1994 binary)
Root relaxation presolved: 8659 rows, 8319 columns, 41712 nonzeros

Root barrier log...

Ordering time: 0.01s

Barrier statistics:
AA' NZ   : 0.906e+04
Factor NZ : 3.210e+05 (roughly 10 MB of memory)
Factor Ops: 1.805e+07 (less than 1 second per iteration)
Threads  : 8
```

CPU model: AMD EPYC 7643 48-Core Processor, instruction set [SSE2|AVX|AVX2]
Thread count: 48 physical cores, 96 logical processors, using up to 8 threads

Optimize a model with 92155 rows, 50135 columns and 262513 nonzeros

Model fingerprint: 0x81f9d80d

Variable types: 29975 continuous, 20160 integer (20160 binary)

Coefficient statistics: $1 \cdot 10^{-8}$ $5 \cdot 10^5$ $Ax \leq b$

Matrix range(A) [1e+02, 5e+02]

Objective range(c) [9e+02, 2e+03]

Bounds range [0e+00, 0e+00]

RHS range(b) [1e+00, 1e+08] \rightarrow suspicious big-M

Found heuristic solution: objective -814886.9986

Presolve removed 83494 rows and 41814 columns

Presolve time: 0.35s

Presolved: 8661 rows, 8321 columns, 41716 nonzeros

Variable types: 6327 continuous, 1994 integer (1994 binary)

Root relaxation presolved: 8659 rows, 8319 columns, 41712 nonzeros

Root barrier log...

Ordering time: 0.01s

Barrier statistics:

AA' NZ : 8.906e+04

Factor NZ : 3.210e+05 (roughly 10 MB of memory)

Factor Ops : 1.805e+07 (less than 1 second per iteration)

Threads : 8

Iter	Objective		Residual		Compl	Time
	Primal	Dual	Primal	Dual		
0	-8.75349353e+10	3.16759838e+11	3.51e+04	3.61e+03	2.85e+07	0s
1	-5.65473859e+10	4.49136693e+09	4.93e+03	9.37e-11	4.19e+06	0s
2	-4.66624680e+09	2.12223081e+09	4.31e+02	9.91e-11	4.16e+05	0s
3	-5.93650592e+08	5.33710820e+08	5.43e+01	1.23e-10	6.09e+04	0s
4	-1.41554632e+08	1.48462275e+08	1.28e+01	9.00e-11	1.46e+04	0s
5	-6.17438000e+07	5.76147676e+07	5.54e+00	5.02e-11	5.82e+03	0s
6	-2.17774077e+07	1.42426566e+07	1.94e+00	3.62e-11	1.71e+03	0s
7	-8.40739562e+06	4.59748907e+06	7.39e-01	1.43e-11	6.06e+02	0s
8	-6.47436121e+06	3.03782035e+06	5.70e-01	1.36e-11	4.41e+02	0s
9	-4.47422322e+06	2.35107115e+06	3.96e-01	9.09e-12	3.16e+02	0s
10	-3.05997438e+06	1.42584289e+06	2.75e-01	7.15e-12	2.07e+02	0s
11	-2.41341865e+06	1.28100790e+06	2.20e-01	6.84e-12	1.70e+02	0s
12	-2.20664641e+06	1.18989676e+06	2.02e-01	8.35e-12	1.57e+02	0s
13	-1.98828288e+06	1.17144512e+06	1.84e-01	8.32e-12	1.46e+02	0s
14	-1.54978590e+06	8.44878454e+05	1.47e-01	5.74e-12	1.10e+02	0s
15	-1.07826885e+06	6.08789586e+05	1.05e-01	5.13e-12	7.78e+01	0s
16	-5.82641701e+05	3.53376691e+05	6.21e-02	4.97e-12	4.31e+01	0s
17	-3.62968591e+05	1.97842701e+05	4.05e-02	2.79e-12	2.59e+01	0s
18	-1.95442249e+05	1.13637501e+05	2.38e-02	3.14e-12	1.43e+01	1s
19	-7.15981068e+04	6.24467786e+04	1.07e-02	1.81e-12	6.19e+00	1s
20	6.33863689e+03	4.23819434e+04	1.85e-03	2.93e-12	1.66e+00	1s
21	1.05186928e+04	2.81456352e+04	1.21e-03	2.18e-12	8.12e-01	1s
22	1.51678553e+04	2.35368659e+04	4.31e-04	1.96e-12	3.84e-01	1s
23	1.66147354e+04	1.99499477e+04	1.89e-04	1.17e-12	1.53e-01	1s
24	1.69590686e+04	1.90483655e+04	1.32e-04	1.18e-12	9.61e-02	1s
25	1.72897664e+04	1.81986844e+04	7.65e-05	1.05e-12	4.20e-02	1s
26	1.75170934e+04	1.79693432e+04	3.80e-05	1.33e-12	2.09e-02	1s
27	1.76892413e+04	1.77892494e+04	8.51e-06	1.76e-12	4.62e-03	1s
28	1.77322731e+04	1.77461969e+04	1.30e-06	1.82e-12	6.44e-04	1s
29	1.77401344e+04	1.77408199e+04	2.73e-09	2.73e-12	3.12e-05	1s
30	1.77401610e+04	1.77401729e+04	1.38e-10	2.73e-12	5.44e-07	1s
31	1.77401626e+04	1.77401627e+04	1.14e-10	2.73e-12	2.03e-09	1s

Barrier solved model in 31 iterations and 0.58 seconds (0.58 work units)

Optimal objective 1.77401626e+04

Root relaxation: objective 1.774016e+04, 6363 iterations, 0.32 seconds (0.30 work units)

Nodes		Current Node			Objective Bounds			Work	
Expl	Unexpl	Obj	Depth	IntInf	Incumbent	BestBd	Gap	It/Node	Time
	0	0	17740.1626	0	581 -814887.00	17740.1626	102%	-	2s
H	0	0			-96621.07624	17740.1626	118%	-	2s
	0	0	17707.4841	0	624 -96621.076	17707.4841	118%	-	3s
	0	0	17704.5241	0	630 -96621.076	17704.5241	118%	-	3s
	0	0	17660.0868	0	588 -96621.076	17660.0868	118%	-	3s
	0	0	17660.0868	0	586 -96621.076	17660.0868	118%	-	3s
H	0	0			-29247.14276	17660.0868	160%	-	6s
	0	0	17660.0868	0	123 -29247.143	17660.0868	160%	-	6s
H	0	0			-14563.64994	17660.0868	221%	-	6s
	0	0	17660.0868	0	123 -14563.650	17660.0868	221%	-	6s
	0	0	17660.0868	0	126 -14563.650	17660.0868	221%	-	6s
	0	0	17660.0868	0	125 -14563.650	17660.0868	221%	-	6s
	0	0	17660.0868	0	231 -14563.650	17660.0868	221%	-	7s
	0	0	17660.0868	0	196 -14563.650	17660.0868	221%	-	7s
	0	0	17660.0868	0	186 -14563.650	17660.0868	221%	-	7s
	0	0	17660.0868	0	194 -14563.650	17660.0868	221%	-	7s
	0	0	17660.0868	0	187 -14563.650	17660.0868	221%	-	7s
	0	0	17660.0868	0	188 -14563.650	17660.0868	221%	-	7s
	0	0	17660.0868	0	151 -14563.650	17660.0868	221%	-	7s
	0	0	17660.0868	0	164 -14563.650	17660.0868	221%	-	7s
	0	0	17660.0868	0	164 -14563.650	17660.0868	221%	-	7s
	0	0	17660.0868	0	169 -14563.650	17660.0868	221%	-	8s
	0	0	17660.0868	0	86 -14563.650	17660.0868	221%	-	8s
H	0	0			-7148.918660	17660.0868	347%	-	9s
	0	2	17660.0868	0	69 -7148.9187	17660.0868	347%	-	9s
	3	8	17660.0868	2	143 -7148.9187	17660.0868	347%	1127	10s
H	32	40			-1440.475097	17510.6047	1316%	791	10s
H	32	40			1318.5313397	17510.6047	1228%	791	10s
H	58	62			3330.7868465	17510.6047	426%	523	11s
H	59	62			4276.5071333	17510.6047	309%	515	11s
H	106	107			4276.5071340	17510.6047	309%	430	12s
H	188	195			4276.5073972	17510.6047	309%	294	13s
H	230	230			4789.4576283	17510.6047	266%	276	14s
H	231	230			5845.1096079	17510.6047	200%	275	14s
	265	259	10627.4584	26	105 5845.10961	17510.6047	200%	273	15s
H	725	448			9351.5655274	17510.6047	87.2%	286	22s
H	726	448			11089.017488	17510.6047	57.9%	286	22s
H	761	433			11089.018283	16804.7463	51.5%	282	22s
	1231	626	11731.0990	59	72 11089.0183	16596.9487	49.7%	264	25s
H	1507	738			11345.897832	16596.9487	46.3%	262	29s
	1508	737	11596.0990	45	86 11345.8978	16596.9487	46.3%	262	30s
	1516	742	11774.9584	71	116 11345.8978	16596.9487	46.3%	260	36s
	1717	776	13460.5704	28	88 11345.8978	16596.9487	46.3%	264	40s
H	1913	746			11345.898980	16596.9487	46.3%	248	42s
	2238	844	12517.9584	32	68 11345.8990	16596.9487	46.3%	233	45s
H	2469	839			11761.920659	16596.9487	41.1%	225	46s
	3079	1054	12020.3437	109	106 11761.9207	16596.9487	41.1%	205	54s
H	3082	1019			11761.921451	16596.9487	41.1%	205	54s
	3092	1024	infeasible	110	11761.9215	16596.9487	41.1%	206	55s
H	3094	810			11774.958412	16596.9487	41.0%	206	55s
	4236	1639	11868.5460	247	101 11774.9584	16596.9487	41.0%	190	60s
H	4304	1639			11774.959215	16596.9487	41.0%	189	60s
	5530	2609	12071.6251	35	104 11774.9592	16058.5000	36.4%	179	65s
	6796	3647	11909.9584	47	100 11774.9592	15531.8030	31.9%	173	70s
H	6903	715			12071.625079	15531.8030	28.7%	174	70s

H 8102	587				12071.627454	14885.4915	23.3%	163	73s
8501	692	12198.6527	36	99	12071.6275	14233.3151	17.9%	161	75s
9402	1180	12821.8972	83	93	12071.6275	13562.7084	12.4%	164	80s
10476	1809	12821.8972	58	91	12071.6275	12824.4084	6.24%	168	86s
11568	2273	infeasible	234		12071.6275	12824.4084	6.24%	172	91s
12694	2779	12574.1774	50	122	12071.6275	12824.4084	6.24%	179	97s
14132	3863	12595.1316	146	79	12071.6275	12824.4084	6.24%	178	102s
14684	4467	12087.5381	225	124	12071.6275	12824.4084	6.24%	176	105s
16345	5169	12821.8972	90	71	12071.6275	12824.4084	6.24%	175	110s
17768	6323	12263.9173	90	94	12071.6275	12824.4084	6.24%	175	115s
19614	7623	cutoff	58		12071.6275	12824.4084	6.24%	173	122s
20605	8210	12294.5174	221	113	12071.6275	12824.4084	6.24%	172	125s
21573	8219	12821.8972	55	89	12071.6275	12824.4084	6.24%	173	130s
22590	9644	12222.1346	146	97	12071.6275	12824.4084	6.24%	172	137s
23664	10334	12821.8972	89	76	12071.6275	12824.4084	6.24%	172	141s
24605	11053	12821.8972	122	102	12071.6275	12824.4084	6.24%	172	145s
26937	12674	12558.4932	97	95	12071.6275	12824.4084	6.24%	173	153s
28404	13642	12481.7713	108	104	12071.6275	12824.4084	6.24%	171	157s
29823	14798	12807.5866	77	102	12071.6275	12824.4084	6.24%	172	207s
31514	15733	12821.8972	134	61	12071.6275	12824.4084	6.24%	171	211s
*32476	15441		268		12074.383586	12824.4084	6.21%	169	211s
32750	16295	12081.3655	351	78	12074.3836	12824.4084	6.21%	169	215s
H34111	15060				12101.373966	12824.4084	5.97%	169	218s
H34454	1258				12821.897187	12824.4084	0.02%	169	218s
35062	843	cutoff	99		12821.8972	12824.4084	0.02%	169	222s
H35069	843				12821.898429	12824.4084	0.02%	169	222s
H35217	843				12821.899229	12824.4084	0.02%	169	222s
35714	768	infeasible	79		12821.8992	12824.4084	0.02%	170	226s
36432	654	12824.4084	92	109	12821.8992	12824.4084	0.02%	173	230s
37040	477	infeasible	81		12821.8992	12824.4084	0.02%	175	235s
38605	305	12824.4084	96	96	12821.8992	12824.4084	0.02%	180	244s
39219	229	cutoff	90		12821.8992	12824.4084	0.02%	183	249s
39947	234	infeasible	49		12821.8992	12824.4084	0.02%	185	253s
40719	249	12824.4084	69	99	12821.8992	12824.4084	0.02%	187	258s
41505	310	infeasible	59		12821.8992	12824.4084	0.02%	188	262s
42399	490	infeasible	96		12821.8992	12824.4084	0.02%	190	267s
43465	584	12824.4084	69	95	12821.8992	12824.4084	0.02%	192	273s
H43852	584				12821.900927	12824.4084	0.02%	192	273s
44046	678	infeasible	77		12821.9009	12824.4084	0.02%	193	279s
45186	811	12824.4084	79	96	12821.9009	12824.4084	0.02%	195	284s
46149	1010	cutoff	86		12821.9009	12824.4084	0.02%	197	291s
47374	1210	12824.4084	70	75	12821.9009	12824.4084	0.02%	200	298s
48794	1435	12824.4084	68	73	12821.9009	12824.4084	0.02%	202	305s
50126	1471	12824.4084	83	107	12821.9009	12824.4084	0.02%	204	312s
51105	1476	cutoff	71		12821.9009	12824.4084	0.02%	206	320s
52398	1421	12824.4084	77	109	12821.9009	12824.4084	0.02%	208	327s
53623	1407	12824.4084	66	97	12821.9009	12824.4084	0.02%	211	335s
54938	1214	12824.4084	90	68	12821.9009	12824.4084	0.02%	213	344s
56249	1057	cutoff	93		12821.9009	12824.4084	0.02%	215	353s
57566	910	12824.4084	81	120	12821.9009	12824.4084	0.02%	218	361s
58936	689	infeasible	76		12821.9009	12824.4084	0.02%	219	371s
60407	532	cutoff	89		12821.9009	12824.4084	0.02%	221	379s
61744	403	12824.4084	108	123	12821.9009	12824.4084	0.02%	223	387s
62935	320	cutoff	93		12821.9009	12824.4084	0.02%	224	394s
63996	276	12824.4084	103	109	12821.9009	12824.4084	0.02%	224	401s
64860	200	cutoff	97		12821.9009	12824.4084	0.02%	225	407s
65647	190	cutoff	65		12821.9009	12824.4084	0.02%	226	412s
66448	156	12824.4084	72	85	12821.9009	12824.4084	0.02%	227	417s
67059	169	12824.4084	84	97	12821.9009	12824.4084	0.02%	227	423s
67792	67	12824.4084	88	92	12821.9009	12824.4084	0.02%	228	428s
68370	16	cutoff	106		12821.9009	12824.4084	0.02%	228	432s

Cutting planes:

Gomory: 4

Projected implied bound: 3
MIR: 85
Flow cover: 8
RLT: 1

Explored 68724 nodes (15775565 simplex iterations) in 434.82 seconds (799.87 work units)
Thread count was 8 (of 96 available processors)

Solution count 10: 12821.9 12821.9 12821.9 ... 11775

Optimal solution found (tolerance 1.00e-04)
Best objective 1.282190092696e+04, best bound 1.282190092696e+04, gap 0.0000%

Typical causes of numerical issues

- Scale: too large of a range of numerical coefficients. Setting parameter `NumericFocus` can help but should be avoided. Ideally, this model should be reformulated.

```
Matrix range    [9e-06, 6e+01]
Objective range [1e+00, 1e+00]
Bounds range    [1e-06, 1e+03]
RHS range       [0e+00, 0e+00]
```

- Rounding of numerical coefficients: Ex: Don't write $1/3$ as 0.333 , if possible multiply constraint with 3
- wrong answer with big-M

```
▸  $y \leq 1000000 \cdot x$ 
   $x$  binary
   $y \geq 0$ 
```

- With default value of `IntFeasTol` ($1e-5$):
 - $x = 0.00000999999$, $y = 9.9999$ is integer feasible!
 - y can be positive without forcing x to 1

Check log file for warnings

- Warning during optimization

```
Warning: 1 variable dropped from basis  
Warning: switch to quad precision  
Warning: Markowitz tolerance tightened to 0.0625
```

- Warnings after optimization is finished

```
Warning: max integrality violation (5.0000e-05) exceeds  
tolerance  
Warning: max SOS violation (9.9353e-05) exceeds tolerance  
Warning: max constraint violation (2.5435e-05) exceeds  
tolerance  
Warning: max bound violation (2.5435e-05) exceeds tolerance  
(possibly due to large matrix coefficient range)  
Numerical trouble encountered
```

References

- CPLEX user manual <https://www.ibm.com/docs/en/icos/22.1.1?topic=optimizers-users-manual-cplex>
- Gurobi Documentation:
<https://www.gurobi.com/documentation/>
- Gurobi Youtube channel
<https://www.youtube.com/@GurobiVideos/featured>
- Paul Rubin, User Cuts versus Lazy Constraints
<https://orinanobworld.blogspot.com/2012/08/user-cuts-versus-lazy-constraints.html>
- Sonja Mars, Modeling: Best Practices & Techniques.
<https://assets.gurobi.com/pdfs/user-events/2016-frankfurt/Best-Practices.pdf>